



CYOK

Choose Your Own Kubernetes

For Local Development



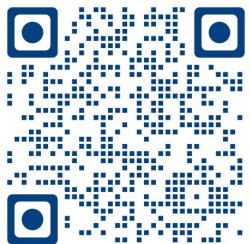
#ospocon @gkarthiks

About me



Karthikeyan Govindaraj
OpenSource Enthusiast | Writer | CNCF Speaker

Developer Evangelist @ BlackRock



<https://bio.link/gkarthiks>

<https://gkarthiks.medium.com>



Concepts

Single Node Cluster: A complete Kubernetes services viz., kube-apiserver, etcd, kube-scheduler, kube-controller-manager running in a single instance of a machine

Multi Node Cluster(HA): A Kubernetes cluster which has multiple worker nodes for workload scheduling, providing HA for the deployed service

HA Control-Plane Cluster: A Kubernetes cluster with multiple worker nodes for deployed workload HA and multiple control-plane node for the Kubernetes cluster HA itself



microK8s



Creates a single/multi node cluster



Built-in container registry as add-on



Start Kubernetes

```
microk8s install  
microk8s status --wait-ready  
microk8s enable dashboard dns registry  
microk8s kubectl get all -A  
alias kubectl="microk8s kubectl"
```

```
ubuntu@microk8s-vm-01:~$ microk8s status --wait-ready  
microk8s is running  
high-availability: no  
  datastore master nodes: 127.0.0.1:19001  
  datastore standby nodes: none  
addons:  
  enabled:  
    dashboard      # The Kubernetes dashboard  
    dns            # CoreDNS  
    ha-cluster     # Configure high availability on the current node  
    metrics-server # K8s Metrics Server for API access to service metrics  
    registry       # Private image registry exposed on localhost:32000  
    storage        # Storage class; allocates storage from host directory  
  disabled:  
    ambassador    # Ambassador API Gateway and Ingress  
    cilium        # SDN, fast with full network policy  
    fluentd      # Elasticsearch-Fluentd-Kibana logging and monitoring  
    gpu           # Automatic enablement of Nvidia CUDA  
    helm          # Helm 2 - the package manager for Kubernetes  
    helm3         # Helm 3 - Kubernetes package manager  
    host-access   # Allow Pods connecting to Host services smoothly  
    ingress       # Ingress controller for external access  
    istio         # Core Istio service mesh services  
    jaeger        # Kubernetes Jaeger operator with its simple config  
    keda          # Kubernetes-based Event Driven Autoscaling  
    knative       # The Knative framework on Kubernetes.  
    kubeflow      # Kubeflow for easy ML deployments  
    linkerd      # Linkerd is a service mesh for Kubernetes and other frameworks  
    metallb     # Loadbalancer for your Kubernetes cluster  
    multus       # Multus CNI enables attaching multiple network interfaces to pods  
    openebs      # OpenEBS is the open-source storage solution for Kubernetes  
    openfaas     # Openfaas serverless framework  
    portainer    # Portainer UI for your Kubernetes cluster  
    prometheus  # Prometheus operator for monitoring and logging  
    rbac         # Role-Based Access Control for authorisation  
    traefik      # traefik Ingress controller for external access
```





Easy multi-node setup



multipass can be used for managing VMs

```
multipass launch --name microk8s-vm-01 --mem 4G --disk 40G  
multipass launch --name microk8s-02 --mem 4G --disk 40G  
multipass launch --name microk8s-03 --mem 4G --disk 40G
```

```
Tue Sep 07@03:46:42 ~/Desktop * > multipass ls
```

Name	State	IPv4	Image
microk8s-02	Running	10.1.63.128	Ubuntu 20.04 LTS
microk8s-03	Running	10.1.62.192	Ubuntu 20.04 LTS
microk8s-vm-01	Running	10.1.17.0	Ubuntu 20.04 LTS





Exec into VM via `multipass shell <control-plane vm-name>`



sudo snap install microk8s in all VMs



Create join token - `microk8s add-node`

```
ubuntu@microk8s-vm-01:~$ microk8s add-node
From the node you wish to join to this cluster, run the following:
microk8s join 10.0.2.15:25000/d6ee805bde2477d4792fee86490e531c/4d8393ca856b

If the node you are adding is not reachable through the default interface you can use one of the following:
microk8s join 10.0.2.15:25000/d6ee805bde2477d4792fee86490e531c/4d8393ca856b
```



with the `add-node` command

enables HA

creates join token

adds the node to the cluster



Run the command in other VMs - `microk8s join` to join the cluster



Highly Available Cluster is ready



enable dashboard:

"microk8s enable dashboard"



setup kubeconfig:

"multipass exec microk8s-vm-01 -- sudo /snap/bin/microk8s.config > kubeconfig"



run from local:

`kubectl --kubeconfig=kubeconfig get all --all-namespaces`



proxy local:

`kubectl --kubeconfig=kubeconfig proxy`



No need of docker daemon to run locally



Can run in all OS viz., Win, Mac, Linux



Built-in container registry for Images



Well suited for CI environment



Supports Windows containers



kind



Creates a single/multi node cluster



Follows the concept of **Kubernetes In Docker**



Owned by k8s SIG



Working with cluster

kind create cluster

```
Tue Sep 07@05:26:53 ~/Desktop * > kind create cluster
Creating cluster "kind" ...
✓ Ensuring node image (kindest/node:v1.21.1) 
✓ Preparing nodes 
✓ Writing configuration 
✓ Starting control-plane 
✓ Installing CNI 
✓ Installing StorageClass 
Set kubectl context to "kind-kind"
You can now use your cluster with:

kubectl cluster-info --context kind-kind

Have a question, bug, or feature request? Let us know! https://kind.sigs.k8s.io/#community 😊
Tue Sep 07@05:27:30 ~/Desktop *
```





Create a custom
build of K8s:
“kind build node-image”



Create & Run
custom built k8s cluster
“kind create cluster --image ”



supports Linux,
macOS and Windows
(with docker daemon)



Images can be loaded from local build



Allows locally built image to directly run inside K8s cluster

```
1 kind: Cluster
2 apiVersion: kind.x-k8s.io/v1alpha4
3 nodes:
4 - role: control-plane
5 - role: control-plane
6 - role: control-plane
7 - role: worker
8 - role: worker
9 - role: worker|
```

Supports kind configuration file

Creates multi-node and control-plane HA cluster

TUE SEP 07@10:12:36	~/Desktop *	> kubectl get no	NAME	STATUS	ROLES	AGE	VERSION
			cp-ha-control-plane	Ready	control-plane,master	4h1m	v1.21.1
			cp-ha-control-plane2	Ready	control-plane,master	4h1m	v1.21.1
			cp-ha-control-plane3	Ready	control-plane,master	4h	v1.21.1
			cp-ha-worker	Ready	<none>	4h	v1.21.1
			cp-ha-worker2	Ready	<none>	4h	v1.21.1
			cp-ha-worker3	Ready	<none>	4h	v1.21.1





Needs docker/podman on the host



Runs in all OS viz., Win/Mac/Linux provided docker daemon is available.



Additional deployment for lib/registry to run container registry for Images



Well suited for CI environment



Doesn't support Windows containers



k3s



Can be installed directly only on Linux



Needs either a VM in local machine or *k3d*



multipass can be used for VMs

```
Terminal  Shell  Edit  View  Session  Scripts  Profiles  Toolbar  Window  Help
ubuntu@control-plane-k3s:~$ curl -sfL https://get.k3s.io | sh -
[INFO]  Finding release for channel stable
[INFO]  Using v1.21.4+k3s1 as release
[INFO]  Downloading hash https://github.com/k3s-io/k3s/releases/download/v1.21.4+k3s1/sha256sum-amd64.txt
[INFO]  Downloading binary https://github.com/k3s-io/k3s/releases/download/v1.21.4+k3s1/k3s
[INFO]  Verifying binary download
[INFO]  Installing k3s to /usr/local/bin/k3s
[INFO]  Creating /usr/local/bin/kubectl symlink to k3s
[INFO]  Creating /usr/local/bin/crictl symlink to k3s
[INFO]  Creating /usr/local/bin/ctr symlink to k3s
[INFO]  Creating killall script /usr/local/bin/k3s-killall.sh
[INFO]  Creating uninstall script /usr/local/bin/k3s-uninstall.sh
[INFO]  env: Creating environment file /etc/systemd/system/k3s.service.env
[INFO]  systemd: Creating service file /etc/systemd/system/k3s.service
[INFO]  systemd: Enabling k3s unit
Created symlink /etc/systemd/system/multi-user.target.wants/k3s.service → /etc/systemd/system/k3s.service.
[INFO]  systemd: Starting k3s
```



<img



K3s can also be run in HA mode



Follow same process to create multiple VMs



Get the control-plane token from “`/var/lib/rancher/k3s/server/node-token`”
& control-plane IP address



Set the token to the env variable “`K3S_TOKEN`”



Install k3s via script with token(automatically picked) and IP arguments as
`curl -sfL https://get.k3s.io | k3S_URL=https://<control-plane-IP> sh -`



k3d - Alternative to VM installations



lightweight wrapper to run k3s in docker

```
Tue Sep 07@13:07:08 ~/Desktop * > k3d cluster create mycluster
INFO[0000] Prep: Network
INFO[0000] Created network 'k3d-mycluster' (feb31d765dc2c19faee37dd706d8bc587e1b94fe82abee3bf494493a8f918789)
INFO[0000] Created volume 'k3d-mycluster-images'
INFO[0001] Creating node 'k3d-mycluster-server-0'
INFO[0006] Pulling image 'docker.io/rancher/k3s:v1.21.3-k3s1'
INFO[0048] Creating LoadBalancer 'k3d-mycluster-serverlb'
INFO[0054] Pulling image 'docker.io/rancher/k3d-proxy:4.4.8'
INFO[0075] Starting cluster 'mycluster'
INFO[0075] Starting servers...
INFO[0075] Starting Node 'k3d-mycluster-server-0'
INFO[0081] Starting agents...
INFO[0081] Starting helpers...
INFO[0081] Starting Node 'k3d-mycluster-serverlb'
INFO[0081] (Optional) Trying to get IP of the docker host and inject it into the cluster as 'host.k3d.internal' for
WARN[0084] Failed to patch CoreDNS ConfigMap to include entry '192.168.65.2 host.k3d.internal': Exec process in node
INFO[0084] Successfully added host record to /etc/hosts in 2/2 nodes
INFO[0084] Cluster 'mycluster' created successfully!
INFO[0084] --kubeconfig-update-default=false --> sets --kubeconfig-switch-context=false
INFO[0084] You can now use it like this:
kubectl config use-context k3d-mycluster
kubectl cluster-info
```





Supports multi node cluster as well

k3d cluster create multiserver --servers 3

```
Tue Sep 07@13:15:13 ~./Desktop * > k3d cluster create multiserver --servers 3
INFO[0000] Prep: Network
INFO[0000] Created network 'k3d-multiserver' (bd6c945ded60338c5154cdcb7276f8928b1444092c45a413937f408862d2e22c)
INFO[0000] Created volume 'k3d-multiserver-images'
INFO[0000] Creating initializing server node
INFO[0000] Creating node 'k3d-multiserver-server-0'
INFO[0001] Creating node 'k3d-multiserver-server-1'
INFO[0002] Creating node 'k3d-multiserver-server-2'
INFO[0002] Creating LoadBalancer 'k3d-multiserver-serverlb'
INFO[0002] Starting cluster 'multiserver'
INFO[0002] Starting the initializing server...
INFO[0002] Starting Node 'k3d-multiserver-server-0'
INFO[0003] Starting servers...
INFO[0003] Starting Node 'k3d-multiserver-server-1'
INFO[0025] Starting Node 'k3d-multiserver-server-2'
INFO[0042] Starting agents...
INFO[0042] Starting helpers...
INFO[0042] Starting Node 'k3d-multiserver-serverlb'
INFO[0046] (Optional) Trying to get IP of the docker host and inject it into the cluster as 'host.k3d.internal' for easy access
INFO[0050] Successfully added host record to /etc/hosts in 4/4 nodes and to the CoreDNS ConfigMap
INFO[0050] Cluster 'multiserver' created successfully!
INFO[0050] --kubeconfig-update-default=false --> sets --kubeconfig-switch-context=false
INFO[0050] You can now use it like this:
kubectl config use-context k3d-multiserver
kubectl cluster-info
```

Auto clustering in multi-server mode

```
Tue Sep 07@13:31:47 ~/Desktop * > kubectl get no
NAME                STATUS    ROLES      AGE     VERSION
k3d-multiserver-server-0  Ready    control-plane,etcd,master  15m    v1.21.3+k3s1
k3d-multiserver-server-1  Ready    control-plane,etcd,master  15m    v1.21.3+k3s1
k3d-multiserver-server-2  Ready    control-plane,etcd,master  14m    v1.21.3+k3s1
```



multiserver control-plane cluster



k3d cluster create mycluster --registry-create



Compatible only with Linux (windows)



Needs multipass/vagrant in local (Linux VMs)



Also run with k3d



Needs docker on the host (for k3d mode)



Additional deployment for lib/registry to run
container registry for Images (native k3s)



Concepts

Virtual Cluster: A complete Kubernetes Cluster spun up inside an existing Kubernetes Cluster

Host Cluster: An actual Kubernetes Cluster that hosts the virtual kubernetes clusters

Self Servicing K8s Cluster: A developer can login and create their own Kubernetes Cluster on demand



vCluster



Doesn't create any cluster in local machine



Creates a cluster inside a running K8s cluster



Follows the *Kubernetes-IN-Kubernetes* concept



Developers works with remote cluster with thin cli



Developers can be (virtual)cluster admin

vCluster



CLI works with all OS viz., Mac/Linux/Windows



No container registry is available, as this will purely creates a Kubernetes cluster only



Namespaces and resources created within the virtual cluster will be encapsulated within the host namespace where vCluster runs



Creates an impartial cloud-native developer experience across all OS Developers



vCluster



Creation of vCluster doesn't need admin perms

```
vcluster create my-vcluster -n host-namespace  
vcluster connect my-vcluster -n host-namespace  
export KUBECONFIG=./kubeconfig.yaml  
kubectl get ns  
kubectl get pods -n kube-system
```



From the host cluster, all the services are sync'd and pods are scheduled from *syncer* a scheduler handler





A host cluster is mandatory



Provides a remote development environment



Need to deploy our own container registry, as this creates a pure K8s cluster only



Impartial cloud-native developer experience for developers from all OS and machines viz., Mac/Win/Linux and VDIs



